

Class09

Kate Zhou(PID: A17373286)

Table of contents

Data Import	1
Exploratory Analysis	2
Winpercent Vs Pricepercent	13
5. Correlation Structure	15
Principal Component Analysis(PCA)	17

Today we will be taking a small step back to look at some data we can taste and explore the correlation structure and principal components of some Halloween candy.

Data Import

```
candy_file <- "candy-data.txt"
candy = read.csv(candy_file, row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer
100 Grand	1	0	1	0	0	1
3 Musketeers	1	0	0	0	1	0
One dime	0	0	0	0	0	0
One quarter	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0
	hard bar	pluribus	sugarpercent	pricepercent	winpercent	
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650

Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q1. How many different candy types are in this dataset?

```
dim(candy)
```

```
[1] 85 12
```

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

Q3. What is your favorite candy in the dataset and what is its winpercent value?

```
candy["Twix", ]$winpercent
```

```
[1] 81.64291
```

Q4. What is the winpercent value for “Kit Kat”?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

Exploratory Analysis

We can use the **skimr** package to get a quick overview of a given dataset. This can be useful for the first time you encounter a new dataset.

```
skimr::skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency: numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?


It looks like the last column `candy$winpercent` is on a different scale to all others.

Q7. What do you think a zero and one represent for the `candy$chocolate` column?

0 means False (isn't chocolate), 1 means True (is chocolate)


Q8. Plot a histogram of winpercent values

```
hist(candy$winpercent, breaks=10)
```



Q9. Is the distribution of winpercent values symmetrical?

```
library(ggplot2)
ggplot(candy) +
  aes(winpercent) +
  geom_histogram(bins=10, fill="lightblue", col="skyblue")
```



No symmetric

Q10. Is the center of the distribution above or below 50%?

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

The center of distribution (Median) is below 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Chocolate is higher ranked than fruity candy

```
choc.candy <- candy[candy$chocolate==1,]
choc.win <- choc.candy$winpercent
fru.candy <- candy[candy$fruity==1,]
fru.win <- fru.candy$winpercent
mean(choc.win)
```

```
[1] 60.92153
```

```
mean(fru.win)
```

```
[1] 44.11974
```

Q12. Is this difference statistically significant?

```
ans <- t.test(choc.win, fru.win)
ans
```

Welch Two Sample t-test

```
data:  choc.win and fru.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

Yes, with a P-value 2.8713778×10^{-8} .

Q13. What are the five least liked candy types in this set?

There are two related runctons that can help here, one is the classic `sort()` and `order()`

```
x <- c (5, 10, 1, 4)
sort(x)
```

```
[1] 1 4 5 10
```

```
x <- c (5, 10, 1, 4)
order(x)
```

```
[1] 3 4 1 2
```

```
inds <- order(candy$winpercent, decreasing = 1)
head(candy[inds,])
```

	chocolate	fruity	caramel	peanut	almond	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1
Reese's pieces	1	0	0		1	0

	crisped	rice	wafer	hard bar	pluribus	sugar	percent
Reese's Peanut Butter cup		0	0	0		0	0.720
Reese's Miniatures		0	0	0		0	0.034
Twix		1	0	1		0	0.546
Kit Kat		1	0	1		0	0.313
Snickers		0	0	1		0	0.546
Reese's pieces		0	0	0		1	0.406

	price	percent	win	percent
Reese's Peanut Butter cup	0.651		84.18029	
Reese's Miniatures	0.279		81.86626	
Twix	0.906		81.64291	
Kit Kat	0.511		76.76860	
Snickers	0.651		76.67378	
Reese's pieces	0.651		73.43499	

```
dec <- order(candy$winpercent)
candy[dec[1:5],]
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0

	crisped	rice	wafer	hard bar	pluribus	sugar	percent	price	percent
Nik L Nip		0	0	0		1	0.197		0.976
Boston Baked Beans		0	0	0		1	0.313		0.511
Chiclets		0	0	0		1	0.046		0.325
Super Bubble		0	0	0		0	0.162		0.116
Jawbusters		0	1	0		1	0.093		0.511

	win	percent
Nik L Nip	22.44534	
Boston Baked Beans	23.41782	
Chiclets	24.52499	
Super Bubble	27.30386	

Jawbusters

28.12744

Q14. What are the top 5 all time favorite candy types out of this set?

```
candy[inds[1:5],]
```

	chocolate	fruity	caramel	peanut	almondy	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1

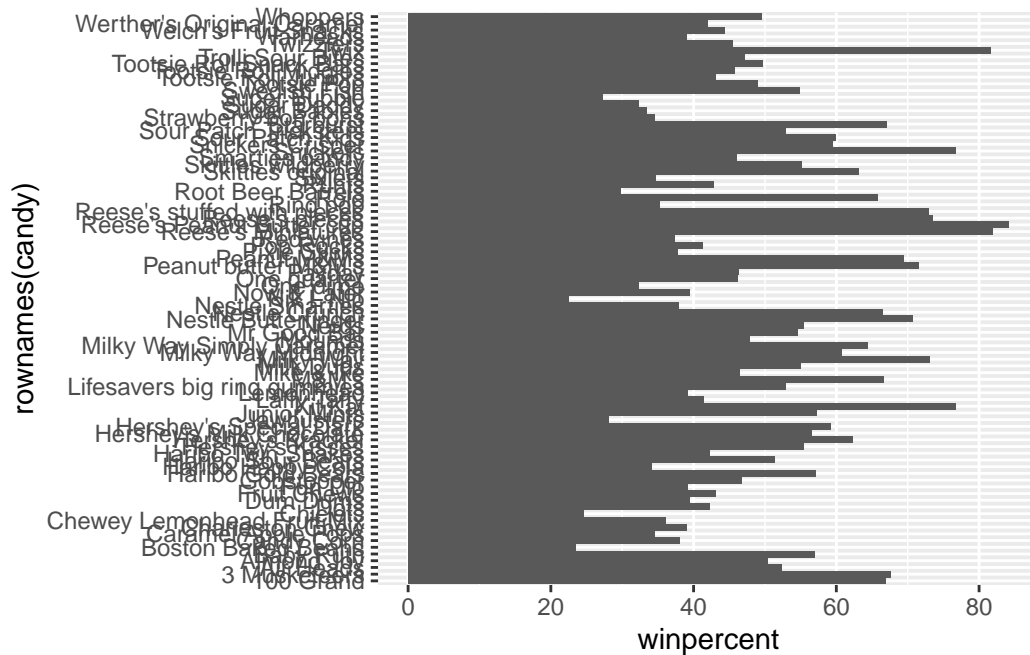
	crisped	rice	wafer	hard	bar	pluribus	sugar	percent
Reese's Peanut Butter cup		0	0	0		0		0.720
Reese's Miniatures		0	0	0		0		0.034
Twix		1	0	1		0		0.546
Kit Kat		1	0	1		0		0.313
Snickers		0	0	1		0		0.546

	price	percent	winpercent
Reese's Peanut Butter cup	0.651		84.18029
Reese's Miniatures	0.279		81.86626
Twix	0.906		81.64291
Kit Kat	0.511		76.76860
Snickers	0.651		76.67378

Q15. Make a first barplot of candy ranking based on winpercent values.


Make a bar plot with ggplot and order it by winpercent values

```
ggplot(candy) +  
  aes(winpercent, rownames(candy)) +  
  geom_col()
```

Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by winpercent?

```
ggplot(candy) +
  aes(x=winpercent,
      y=reorder(rownames(candy), winpercent),
      fill=chocolate) +
  geom_col()
```



Here we want a custom color vector to color each bar the way we want - with **chocolate** and **fruity** candy together with whether it is a **bar** or not

```
mycols <- rep("grey", nrow(candy))
mycols[as.logical(candy$chocolate)] <- "chocolate"
mycols[as.logical(candy$fruity)] <- "pink"
mycols[as.logical(candy$bar)] <- "brown"

ggplot(candy) +
  aes(x=winpercent,
      y=reorder(rownames(candy), winpercent)) +
  geom_col(fill=mycols)
```





Figure 1: Barplot

Q17. What is the worst ranked chocolate candy?

```
worst <- candy[dec,]  
worst_choc <- worst[worst$chocolate == 1,]  
worst_choc[1,]
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
Sixlets	1	0	0	0	0	0	0

	bar	pluribus	sugarpercent	pricepercent	winpercent
Sixlets	0	1	0.22	0.081	34.722

Q18. What is the best ranked fruity candy?

```
best <- candy[inds,]  
best_fruit <- best[best$fruity == 1,]  
best_fruit[1,]
```


	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
Starburst	0	1	0	0	0	0	0

	bar	pluribus	sugarpercent	pricepercent	winpercent
Starburst	0	1	0.151	0.22	67.03763

Winpercent Vs Pricepercent

```
library(ggrepel)  
  
mycols <- rep("grey", nrow(candy))  
mycols[as.logical(candy$chocolate)] <- "chocolate"  
mycols[as.logical(candy$fruity)] <- "red"  
mycols[as.logical(candy$bar)] <- "brown"  
  
# How about a plot of price vs win  
ggplot(candy) +  
  aes(winpercent, pricepercent, label=rownames(candy)) +  
  geom_point(col=mycols) +  
  geom_text_repel(col=mycols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Reese's Miniatures.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
expensive5 <- candy[order(candy$pricepercent, decreasing = 1)[1:5],]
expensive5
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Nik L Nip	0	1	0		0	0
Nestle Smarties	1	0	0		0	0
Ring pop	0	1	0		0	0
Hershey's Krackel	1	0	0		0	0
Hershey's Milk Chocolate	1	0	0		0	0

	crisped	rice	wafer	hard	bar	pluribus	sugarpercent
Nik L Nip			0	0	0	1	0.197
Nestle Smarties			0	0	0	1	0.267
Ring pop			0	1	0	0	0.732
Hershey's Krackel			1	0	1	0	0.430
Hershey's Milk Chocolate			0	0	1	0	0.430

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

```
expensive5[order(expensive5$winpercent)[1],]
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
Nik L Nip	0	1	0	0	0	0	0

	bar	pluribus	sugarpercent	pricepercent	winpercent
Nik L Nip	0	1	0.197	0.976	22.44534

5. Correlation Structure

```
library(corrplot)
```

```
corrplot 0.95 loaded
```


```
cij <- cor(candy)
cij
```

	chocolate	fruity	caramel	peanutyalmondy	nougat
chocolate	1.0000000	-0.74172106	0.24987535	0.37782357	0.25489183
fruity	-0.7417211	1.00000000	-0.33548538	-0.39928014	-0.26936712
caramel	0.2498753	-0.33548538	1.00000000	0.05935614	0.32849280
peanutyalmondy	0.3778236	-0.39928014	0.05935614	1.00000000	0.21311310
nougat	0.2548918	-0.26936712	0.32849280	0.21311310	1.00000000
crispedricewafer	0.3412098	-0.26936712	0.21311310	-0.01764631	-0.08974359
hard	-0.3441769	0.39067750	-0.12235513	-0.20555661	-0.13867505
bar	0.5974211	-0.51506558	0.33396002	0.26041960	0.52297636
pluribus	-0.3396752	0.29972522	-0.26958501	-0.20610932	-0.31033884
sugarpercent	0.1041691	-0.03439296	0.22193335	0.08788927	0.12308135
pricepercent	0.5046754	-0.43096853	0.25432709	0.30915323	0.15319643
winpercent	0.6365167	-0.38093814	0.21341630	0.40619220	0.19937530

	crispedricewafer	hard	bar	pluribus
chocolate	0.34120978	-0.34417691	0.59742114	-0.33967519

fruity	-0.26936712	0.39067750	-0.51506558	0.29972522
caramel	0.21311310	-0.12235513	0.33396002	-0.26958501
peanutyalmondy	-0.01764631	-0.20555661	0.26041960	-0.20610932
nougat	-0.08974359	-0.13867505	0.52297636	-0.31033884
crispedricewafer	1.00000000	-0.13867505	0.42375093	-0.22469338
hard	-0.13867505	1.00000000	-0.26516504	0.01453172
bar	0.42375093	-0.26516504	1.00000000	-0.59340892
pluribus	-0.22469338	0.01453172	-0.59340892	1.00000000
sugarpercent	0.06994969	0.09180975	0.09998516	0.04552282
pricepercent	0.32826539	-0.24436534	0.51840654	-0.22079363
winpercent	0.32467965	-0.31038158	0.42992933	-0.24744787
	sugarpercent	pricepercent	winpercent	
chocolate	0.10416906	0.5046754	0.6365167	
fruity	-0.03439296	-0.4309685	-0.3809381	
caramel	0.22193335	0.2543271	0.2134163	
peanutyalmondy	0.08788927	0.3091532	0.4061922	
nougat	0.12308135	0.1531964	0.1993753	
crispedricewafer	0.06994969	0.3282654	0.3246797	
hard	0.09180975	-0.2443653	-0.3103816	
bar	0.09998516	0.5184065	0.4299293	
pluribus	0.04552282	-0.2207936	-0.2474479	
sugarpercent	1.00000000	0.3297064	0.2291507	
pricepercent	0.32970639	1.0000000	0.3453254	
winpercent	0.22915066	0.3453254	1.0000000	

```
corrplot(cij)
```

Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolate and fruity are negatively correlated

```
round(cij["chocolate","fruity"],2)
```

```
[1] -0.74
```

Q23. Similarly, what two variables are most positively correlated?

```
round(cij["chocolate","winpercent"],2)
```

```
[1] 0.64
```

Principal Component Analysis(PCA)

We need to be sure to scale our input `candy` data before PCA as we have `winpercent` column on a different scale to all others in the dataset.

```
pca <- prcomp(candy, scale=T)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

First main result figure is my “PCA plot”

```
head(pca$x)
```

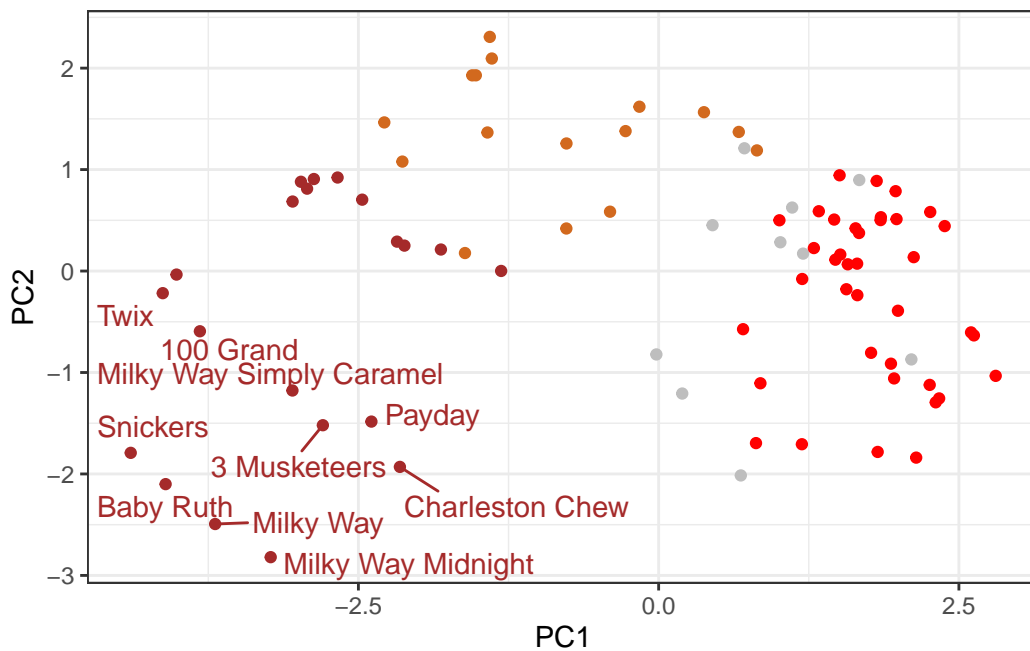
	PC1	PC2	PC3	PC4	PC5
100 Grand	-3.8198617	-0.5935788	-2.1863087	-2.3715957	-0.66236243
3 Musketeers	-2.7960236	-1.5196062	1.4121986	0.6994387	-0.16006665
One dime	1.2025836	0.1718121	2.0607712	-1.2006782	-0.26977985
One quarter	0.4486538	0.4519736	1.4764928	-1.0017714	-0.05093737
Air Heads	0.7028992	-0.5731343	-0.9293893	0.4124566	0.33108524
Almond Joy	-2.4683383	0.7035501	0.8581089	0.5724974	1.43200435

	PC6	PC7	PC8	PC9	PC10
100 Grand	-0.54521840	-0.1434056	0.5772242	0.3791482	-0.15409954
3 Musketeers	0.38258842	2.0215553	-1.5025750	0.0238327	-0.15523907
One dime	0.09495053	-0.7722007	0.1556221	-0.4539890	-0.94378362
One quarter	0.42835404	-0.5908920	0.3591344	-0.6388898	0.09557965
Air Heads	0.18879160	0.9680808	1.0123933	-1.3830122	-0.57372349
Almond Joy	1.02227348	-0.6660460	0.6082613	-0.1390599	-0.31277870

	PC11	PC12
100 Grand	0.1419038	0.06469883
3 Musketeers	-0.5179272	-0.18394717
One dime	-0.5158708	-0.51074779
One quarter	-0.9867194	-0.95827191
Air Heads	-0.5144537	0.04429924
Almond Joy	1.0611487	0.18893471

```
ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(col=mycols) +
  geom_text_repel(max.overlaps = 4, col=mycols) +
  theme_bw()
```

Warning: ggrepel: 75 unlabeled data points (too many overlaps). Consider increasing max.overlaps



The second main PCA result is in the `pca$rotation`, we can plot this to generate a so-called “loadings” plot.


```
head(pca$rotation)
```

	PC1	PC2	PC3	PC4	PC5
chocolate	-0.4019466	0.21404160	0.01601358	-0.016673032	0.06603585
fruity	0.3683883	-0.18304666	-0.13765612	-0.004479829	0.14353533
caramel	-0.2299709	-0.40349894	-0.13294166	-0.024889542	-0.50730150
peanutyalmondy	-0.2407155	0.22446919	0.18272802	0.466784287	0.39993025
nougat	-0.2268102	-0.47016599	0.33970244	0.299581403	-0.18885242
crispedricewafer	-0.2215182	0.09719527	-0.36485542	-0.605594730	0.03465232

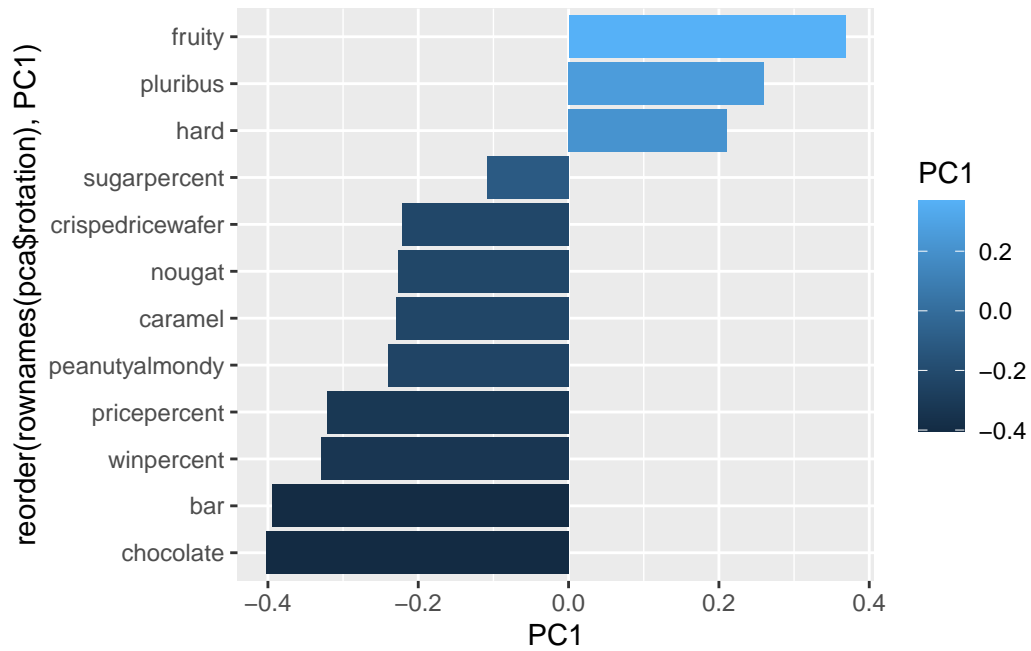
	PC6	PC7	PC8	PC9	PC10
chocolate	-0.09018950	-0.08360642	-0.4908486	-0.151651568	0.10766136
fruity	-0.04266105	0.46147889	0.3980580	-0.001248306	0.36206250
caramel	-0.40346502	-0.44274741	0.2696345	0.019186442	0.22979901
peanutyalmondy	-0.09416259	-0.25710489	0.4577145	0.381068550	-0.14591236
nougat	0.09012643	0.36663902	-0.1879396	0.385278987	0.01132345
crispedricewafer	-0.09007640	0.13077042	0.1356774	0.511634999	-0.26481014

	PC11	PC12
chocolate	0.1004528	0.69784924
fruity	0.1749490	0.50624242
caramel	0.1351582	0.07548984
peanutyalmondy	0.1124428	0.12972756
nougat	-0.3895447	0.09223698
crispedricewafer	-0.2261562	0.11727369

```
ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation)) +
  geom_col()
```



```
#reorder(pca$rotation, winpercent)
ggplot(pca$rotation) +
  aes(PC1, reorder(rownames(pca$rotation), PC1), fill = PC1) +
  geom_col()
```



Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity, pluribus, hard